

JERUSALEM COLLEGE OF ENGINEERING

(AN AUTONOMOUS INSTITUTION)

Approved by AICTE, New Delhi | Affiliated to Anna University, Chennai | Accredited by NBA & NAAC with 'A' Grade
Department of Artificial Intelligence and Machine Learning

Degree / Branch: B.E. CSE (AI & ML)	Course Code / Name: JAL2007 / Storage Technologies
Year / Sem / Regulation: III / V / Regulation 2023	Unit & Question: Unit IV / Part B Question No. 3
Bloom's Taxonomy: K3 (Apply Level)	Course Outcome: CO4 (Backup, Archive & Replication)

QUESTION BANK PART B (16 MARKS):

'Apply data deduplication techniques in backup environments and their working principles. Illustrate how their advantages and limitations affect storage utilization and backup performance.'

MODEL ANSWER INDEX & PAGE STRUCTURE

Se c	Topic Description	Page
1	Introduction to Data Deduplication & Foundational Concepts	Page 1
2	Deduplication Architectural Approaches & Implementation Granularity	Page 2
3	Core Working Principles & Algorithmic Execution Workflow	Page 3
4	Source-Side vs. Target-Side Deduplication & Deployment Topologies	Page 4
5	Impact on Storage Utilization & Backup/Restore Performance Parameters	Page 5
6	Comparative Evaluation, Trade-offs & Enterprise Recommendation Matrix	Page 6

1. INTRODUCTION TO DATA DEDUPLICATION & FOUNDATIONAL CONCEPTS

Data Deduplication is a specialized data compression technique designed to eliminate redundant copies of repeating data across secondary storage and backup infrastructures. Instead of retaining multiple identical byte sequences or file copies across daily full and incremental backup streams, deduplication identifies duplicate blocks, stores only a single physical instance (the 'canonical block'), and replaces all subsequent duplicate occurrences with lightweight cryptographic pointers (metadata references).

In modern enterprise data centers, organizations perform daily, weekly, and monthly backups to fulfill stringent Business Continuity (BC) and Disaster Recovery (DR) requirements. Without deduplication, traditional backup strategies lead to extreme data sprawl, where 80% to 90% of the daily ingested backup data consists of unchanged operating system files, identical application binaries, shared virtual machine templates (VMDK/VHDX), and duplicate email attachments.

Core Objectives of Deduplication in Backup Environments:

- **Storage Capacity Optimization:** Drastically reduces physical storage footprints by achieving compression ratios ranging from 10:1 (90% savings) up to 30:1 (96.7% savings) for highly redundant backup streams.
- **Network Bandwidth Reduction:** Minimizes the volume of payload transferred across Local Area Networks (LAN) and Wide Area Networks (WAN) during remote backup and Disaster Recovery replication.
- **Extended On-Disk Retention:** Enables cost-effective, long-term disk-based retention (Disk-to-Disk / D2D), eliminating total dependency on legacy magnetic tape libraries for historical compliance.
- **Enhanced TCO & Energy Efficiency:** Lowers total cost of ownership (TCO) by reducing raw disk purchasing costs, rack space requirements, power consumption, and data center cooling overheads.

2. DEDUPLICATION ARCHITECTURAL APPROACHES & IMPLEMENTATION GRANULARITY

Deduplication mechanisms are classified based on two fundamental architectural axes: **(1) Implementation Granularity** (File-level vs. Block-level), and **(2) Timing of Execution** (In-line vs. Post-processing). Selecting the appropriate architectural approach dictates CPU utilization, RAM requirements, and overall backup throughput.

2.1 Granularity of Deduplication (File-Level vs. Block-Level)

A. File-Level Deduplication (Single-Instance Storage - SIS):

File-level deduplication operates by comparing full files based on their metadata attributes (file name, path, creation date, file size) or by computing a hash value across the entire file contents. If two files are identical, only one copy is retained, and subsequent files are replaced with short logical pointers.

Limitation: File-level deduplication fails if even a single byte within a 500 MB presentation or database file changes; the entire modified file is treated as brand-new data and written in full, leading to massive redundant storage consumption.

B. Block-Level Deduplication (Chunk-Based Deduplication):

Block-level deduplication breaks incoming file streams or disk images into small logical blocks or 'chunks' (e.g., 4 KB to 128 KB). The system processes each block independently, evaluating its cryptographic hash key against an active index database. If a 10 GB database file experiences a 2 MB change, block-level deduplication stores only the modified 2 MB block while pointing all unchanged blocks to previously saved chunks. Block-level deduplication is divided into two methods:

- **Fixed-Length Chunking:** Divides data into uniform, fixed-size blocks (e.g., exactly 8 KB). While computationally fast and easy to index, fixed-length chunking suffers from the 'boundary shift problem'. Inserting a single byte at the beginning of a file shifts all subsequent block boundaries, rendering the entire file non-matchable.
- **Variable-Length Chunking (Rabin Fingerprinting):** Utilizes rolling hash algorithms to inspect data patterns and dynamically insert chunk boundaries based on specific byte sequences. Even if bytes are inserted or deleted, the boundary shifts locally, leaving all surrounding chunks perfectly identifiable and deduplicable.

2.2 Execution Timing (In-Line vs. Post-Processing Deduplication)

Architectural Factor	In-Line Deduplication	Post-Processing Deduplication
Working Mechanism	Data is deduplicated in RAM/cache <i>before</i> it is written to physical backup disk storage.	Data is written in raw form to a landing disk zone, and deduplicated asynchronously later.
Landing Zone Need	Zero landing space needed. Saves disk space immediately upon ingestion.	Requires large staging disk space equal to 100% of the raw un-deduplicated backup.
CPU & RAM Impact	High real-time CPU/RAM consumption during the active backup window.	No active CPU/RAM load during backup window; heavy CPU usage during off-peak hours.
Backup Speed	Slightly slower ingestion due to real-time hashing and index lookup latency.	Maximum ingestion speed because data is written directly to high-speed staging disk.
Restore Speed	Requires instant rehydration penalty during data recovery.	Ultra-fast recovery if restore occurs before post-processing deduplication finishes.

EXAM KEY INSIGHT: Variable-length in-line block deduplication is the industry standard for enterprise backup appliances (e.g., Dell EMC Data Domain) because it eliminates the need for expensive raw staging disks and minimizes physical storage footprints instantly.

3. CORE WORKING PRINCIPLES & ALGORITHMIC EXECUTION WORKFLOW

The algorithmic engine of a block-level deduplication system relies on cryptographic hash functions, hash index databases, and metadata mapping tables to process incoming backup data streams. The execution workflow follows five sequential operational steps:

Step 1: Ingestion & Chunking:

The incoming backup payload stream is received by the deduplication engine and parsed into discrete data blocks. The system applies fixed-size slicing or variable-length Rabin Fingerprinting algorithms to establish chunk boundaries.

Step 2: Cryptographic Hash Generation:

For every individual data block, the engine computes a unique mathematical signature (fingerprint) using cryptographic hash algorithms such as SHA-256 or MD5. The probability of two different blocks generating the identical SHA-256 hash (known as a *hash collision*) is less than 1 in 10^{77} , making it cryptographically unique.

Step 3: Hash Index Lookup & Matching:

The engine compares the newly generated hash key against an in-memory active Hash Index Table. If the hash key already exists in the index table, a duplicate block is detected. If the key is absent, the block is identified as a unique 'canonical block'.

Step 4: Unique Storage & Reference Pointer Assignment:

- *For Unique Blocks:* The unique block is written to physical disk storage, and its hash key is registered in the Hash Index Table along with its disk location address.
- *For Duplicate Blocks:* The duplicate block is immediately discarded from memory. The deduplication engine generates a lightweight metadata pointer (8–16 bytes) referencing the existing physical block location.

Step 5: Metadata Pointer Map Construction:

The backup software updates the logical file catalog map (the 'rehydration recipe'). This metadata map maintains the exact sequence of original blocks and pointer references, ensuring that the full file structure can be perfectly reconstructed during data restoration.

3.1 Mathematical Principles & Rehydration Formula

DEDUPLICATION MATHEMATICAL FORMULAS:

1. Deduplication Ratio (D_ratio):

$$D_ratio = \text{Total Raw Un-deduplicated Data} / \text{Total Physical Storage Consumed}$$

2. Space Savings Percentage (S_percent):

$$S_percent = (1 - (1 / D_ratio)) * 100\%$$

3. Example Calculation: An organization ingests 100 TB of total backup data over 10 cycles, but stores only 5 TB of physical data on disk.

- $D_ratio = 100 \text{ TB} / 5 \text{ TB} = 20:1$
- $S_percent = (1 - (1 / 20)) * 100\% = 95\% \text{ Storage Savings!}$

3.2 Hash Collision Protection Mechanisms

Because hash collision could result in catastrophic data corruption (where a new unique block is misidentified as a duplicate and discarded), enterprise deduplication appliances employ strict safety controls: **(1) Extended SHA-256 Hash Bit Length**, providing 2^{256} unique address spaces, and **(2) Byte-by-Byte Verification Option**, where the engine performs a full byte comparison if a hash match is flagged before dropping the block.

4. SOURCE-SIDE VS. TARGET-SIDE DEDUPLICATION & DEPLOYMENT TOPOLOGIES

Deduplication can be deployed at different physical locations within the backup infrastructure. The two primary deployment topologies are **Source-Side Deduplication** (executed on the production host server) and **Target-Side Deduplication** (executed on dedicated backup storage appliances).

4.1 Source-Side Deduplication Mechanics

In **Source-Side Deduplication**, a lightweight backup agent software running on the client or application server performs chunking and hashing directly on the client system. The agent communicates with the backup server's central hash index to determine duplicate blocks *before* sending data over the network.

- **Primary Advantage:** Massive network bandwidth savings (up to 90%+ reduction). Only new unique data blocks travel across the LAN/WAN network.
- **Primary Drawback:** Consumes CPU processing power and memory RAM resources from the production server, which can degrade mission-critical database performance during active backup runs.
- **Ideal Use Case:** Remote branch office (ROBO) backups, mobile endpoints, virtual machine environments (e.g., VMware vSphere / Veeam), and low-bandwidth cloud backup over WAN.

4.2 Target-Side Deduplication Mechanics

In **Target-Side Deduplication**, production servers send uncompressed raw backup payloads over the network to a dedicated specialized backup target system (e.g., Dell EMC Data Domain, HPE StoreOnce). The target system receives the full data stream and executes chunking, hashing, deduplication, and compression internally.

- **Primary Advantage:** Offloads 100% of the hashing and CPU processing overhead from production servers. Production systems experience zero overhead.
- **Primary Drawback:** Requires a high-bandwidth Local Area Network (LAN) infrastructure, as raw un-deduplicated data must travel across the network to reach the storage target.
- **Ideal Use Case:** Central enterprise data centers with high-speed 10GbE / Fibre Channel fabrics and heavy database workloads.

Comparison Parameter	Source-Side Deduplication	Target-Side Deduplication
Processing Location	Production Server / Client Agent	Dedicated Backup Storage Appliance
Network Bandwidth Usage	Ultra-low (Transmits unique blocks only)	High (Transmits full raw payload over network)
Host CPU/RAM Impact	High CPU & RAM overhead on host system	Zero impact on production host system
Deployment Cost	Software-based; lower hardware expenditure	Higher initial capital cost for storage target
Scalability Focus	Scales across client computing nodes	Scales central storage target capacity

5. IMPACT ON STORAGE UTILIZATION & BACKUP/RESTORE PERFORMANCE PARAMETERS

While data deduplication delivers transformative storage savings, it significantly alters system performance dynamics across backup ingestion windows, restore execution speeds, and CPU/RAM memory utilization.

5.1 Impact on Storage Utilization & Infrastructure Capacity

Deduplication dramatically compresses secondary storage footprint requirements. For standard enterprise workloads (mix of OS images, virtual machines, file shares, and databases), deduplication achieves average storage reduction ratios between **10:1 (90%)** and **20:1 (95%)**. When combined with traditional LZ4 or gzip compression, cumulative savings can reach **30:1 (96.7%)**.

- **Extended Retention:** Organizations can retain 12 months of disk-based recovery points within the physical footprint that previously stored only 30 days of raw backups.
- **Disaster Recovery Vaulting:** Secondary off-site storage targets require only a fraction of physical disk capacity, lowering cloud storage costs for AWS S3 / Azure Blob repositories.

5.2 Backup Window & Ingestion Throughput Dynamics

A. Source-Side Deduplication Impact: Reduces backup completion times dramatically over slow WAN links because only unique chunks travel over the wire. However, if host CPU is throttled, client-side hashing can create an ingestion bottleneck.

B. In-Line Target Deduplication Impact: Requires ultra-fast hardware controllers equipped with NVMe flash caches and dedicated ASIC/FPGA cryptographic accelerators to maintain high ingestion speeds (e.g., 15–30 TB/hour) without overflowing input buffers.

5.3 Restore Performance Penalty (The Data Rehydration Overhead)

The most critical tradeoff of deduplication is the **Data Rehydration Penalty** during recovery operations. Rehydration is the process of reconstructing a requested logical file by looking up its metadata recipe, fetching hundreds or thousands of fragmented non-contiguous blocks scattered across physical disks, and assembling them back into sequential byte streams.

- **Random Read Overhead:** Because deduplicated blocks are shared globally, restoring a single VM or database requires thousands of random read operations across disk sectors, leading to significant disk head latency on magnetic HDDs.
- **Rehydration CPU Latency:** The storage target must process metadata pointers and uncompress data blocks simultaneously, resulting in slower restore times compared to un-deduplicated sequential raw disk reads.

5.4 Interaction with Data Encryption & Compression

CRITICAL TECHNICAL RULE — ORDER OF OPERATIONS:

Encryption and Compression MUST be performed AFTER Deduplication, never before.

- Cryptographic encryption converts structured data into pseudo-random high-entropy byte patterns. If data is encrypted on the client before reaching the deduplication engine, identical blocks produce completely different ciphertexts, reducing deduplication efficiency to **0%**!
- *Correct Architecture Pipeline:* **[Chunking & Hashing] → [Deduplication Match/Discard] → [LZ4 Compression] → [AES-256 Encryption] → [Disk Storage]**.

6. COMPARATIVE EVALUATION, TRADE-OFFS & ENTERPRISE RECOMMENDATION MATRIX

To synthesize the overall trade-offs of deduplication techniques, the table below provides a comprehensive multi-attribute evaluation across 8 key criteria, followed by quantitative performance chart insights.

Evaluation Criteria	Source-Side Dedup	Target-Side In-Line	Target-Side Post-Proc
Storage Capacity Reduction	Very High (10:1–20:1)	Very High (10:1–30:1)	Very High (10:1–30:1)
Network Bandwidth Impact	Minimal (<10% load)	High (100% raw payload)	High (100% raw payload)
Production Host CPU Overhead	High (15–25% usage)	Zero Impact	Zero Impact
Backup Execution Window	Fast over WAN/LAN	Moderate (Ingestion bound)	Ultra-Fast Ingestion
Data Recovery / Restore Speed	Moderate (Rehydration)	Moderate (Rehydration)	Ultra-Fast (Recent Data)
Landing Zone Disk Requirement	None	None	High (100% raw staging)
Primary Enterprise Use Case	ROBO, Cloud & Endpoint	Core Data Center SAN	High-speed DB Ingestion




Figure 6.1: Quantitative Comparison of Storage Capacity Growth and Backup Execution Windows across Deduplication Techniques.

6.1 Strategic Enterprise Architectural Selection Guidelines

- 1. For Remote Office / Branch Office (ROBO) & Cloud Workloads: Deploy **Source-Side Variable-Length Deduplication** to minimize WAN bandwidth consumption over slow remote links.
- 2. For High-Performance Enterprise Data Centers: Deploy **Target-Side In-Line Deduplication** utilizing NVMe-backed flash storage targets to isolate production hosts from CPU overhead.
- 3. For Mission-Critical Recovery Tiering: Utilize **Hybrid Deduplication with Instant VM Recovery (Virtual Landing Zone)**, retaining the most recent full backup in un-deduplicated raw state for instant recovery, while archiving older retention sets in deduplicated block pools.

Conclusion & Summary: Data deduplication is an essential cornerstone of modern backup architecture. By combining variable-length block chunking, SHA-256 cryptographic hashing, and metadata pointer maps, deduplication reduces secondary storage requirements by over 90%. Understanding the trade-offs between source-side and target-side processing, as well as managing the rehydration penalty during restores, enables storage architects to build resilient, cost-effective, and highly performant enterprise backup solutions.